

# GPU-Based Simulation and Reinforcement Learning Pipeline for Tracked Ground Robots

David Korčák

Prague, January 2025

# Introduction

- Original idea: train an RL policy for MARV or TRADR to control their flippers to minimize bouncing, angular speeds etc.
- What actually happened: neither of the simulators were ready to work with flippers and/or were buggy
- PyTorch simulator was also quite slow
- New objective: develop a PyTorch-based simulator able to handle flipper control, be more numerically stable, bug-free and significantly faster


# Background

- Why not use an existing simulation engine?
- Popular simulation engines like mujoco or brax are optimized for robots with joints and minimal collisions.
- No native support for tracked/wheeled robots and complex terrain collisions.


# Simulation Engine

- Primarily designed for MARV and TRADR
- Simulation based on a pointcloud model of the robot.
- Implemented more elaborate and detail-preserving mesh-to-pointcloud conversion.
- Interactions modeled using spring-damper systems for terrain forces.
- Efficient computation of collision detection and dynamics.


# Mesh to Pointcloud Conversion




# Mesh to Pointcloud Conversion




(a) Full flipper mesh.



(b) Delaunay triangulation



(c) Extracted surface points.




(d) Clustered surface points.

# GPU Optimization


- Leveraging PyTorch for gpu-accelerated matrix operations.
- Conditional branching removed to avoid GPU thread divergence and halting.
- Torch's new compile functionality used for kernel optimization.

# Forward Computational Graph






# Performance Evaluation - CPU (M1 Pro)




# Performance Evaluation - GPU (A100)




# Reinforcement Learning Environment

- Implemented a basic RL framework and environment using TorchRL (native interop with Tensors, devices, datatypes, shapes).
- Observation types: bird-view heightmap camera and 3d lidar pointcloud.
- Random procedural terrain generation.

# Procedural Terrain Generation



(a) Generated terrain with rough preset.



(b) Generated terrain with smooth preset.

# Conclusion

- GPU-friendly simulation engine and basic rl environment developed.
- Future work: refine physics models and complete rl experiments.
- Merge with Aleš and his more advanced joint modeling